

Adaptive Lyapunov-Guided Frequency and Gain Scheduling for PID–MPC Hybrid Control of Nano-UAVs Under Wind and Payload Disturbances

Pappu Meghana¹ and Sakthi Swarrup J²

¹Dept. of AI&DS, IIIT Sri City, Chittoor, India

²Dept. of CSE, IIIT Sri City, Chittoor, India

pappumeghana2006@gmail.com; sakthiswarrup.j@iiits.in

Abstract

This paper presents an adaptive hybrid control architecture for nano-unmanned aerial vehicles (UAVs) that integrates a Lyapunov stability monitor with a Model Predictive Controller (MPC) outer loop and a Proportional-Integral-Derivative (PID) inner loop. The proposed framework addresses two fundamental limitations of fixed-schedule hybrid controllers: (1) the inability to adapt computational resources to current stability demands, and (2) static gain tuning that cannot respond to time-varying disturbances such as wind forces and payload mass changes. A quadratic Lyapunov candidate function $V(\mathbf{x}_e) = \mathbf{x}_e^\top P \mathbf{x}_e$ is constructed from the terminal cost matrix P of the discrete-time Linear Quadratic Regulator (LQR) associated with the MPC formulation. The Lyapunov stability margin $\sigma = -\dot{V}/(V + \varepsilon)$ drives two online mechanisms: (i) an adaptive MPC frequency scheduler that tightens or relaxes the MPC solve rate between 12 Hz and 48 Hz based on stability, and (ii) an adaptive PID gain scheduler that scales proportional, integral, and derivative gains in proportion to the stability deficit and tracking error. Simulations on a Crazyflie 2.x nano-UAV in PyBullet across Hover, Circle, and Figure-8 trajectories under four disturbance conditions demonstrate that the adaptive hybrid controller reduces RMSE by up to 83% on dynamic trajectories compared to fixed-rate MPC, prevents the complete actuator-saturation failure that fixed MPC suffers under combined wind and payload disturbances (RMSE = 2.24 m for fixed MPC versus 0.27 m for PID+MPC on the Circle Wind+Payload trial), and achieves lower mean step time than fixed MPC through Lyapunov-guided frequency relaxation during stable flight phases.

Index Terms

Lyapunov stability, Model Predictive Control, PID control, adaptive gain scheduling, UAV, nano-drone, hybrid control, frequency scheduling, Crazyflie.

I. INTRODUCTION

Nano-unmanned aerial vehicles (UAVs), such as the Crazyflie 2.x platform (mass 27 g), operate at the intersection of severe physical constraints and demanding real-time control requirements. Their low inertia makes them acutely sensitive to external disturbances including wind gusts and payload mass variations, while their embedded computational resources impose strict limits on the complexity of onboard control algorithms. Two classical strategies have been extensively studied for such platforms: Proportional-Integral-Derivative (PID) control, prized for its computational frugality; and Model Predictive Control (MPC), valued for its ability to anticipate future states and enforce constraint satisfaction. Neither approach alone is optimal for nano-UAVs under real-world conditions.

PID controllers, when operating with fixed gains, are tuned for a nominal operating point and exhibit degraded performance when system parameters change. Payload attachment changes the effective mass by more than 55% for the Crazyflie 2.x, directly altering the thrust-to-weight ratio and making nominal gain settings insufficient. Wind disturbances introduce persistent lateral forces that a fixed-gain integrator may track only slowly. MPC, while capable of handling such conditions through its receding-horizon optimisation, incurs significant computational overhead: a convex quadratic programme must be solved at each execution cycle.

A hybrid PID+MPC architecture addresses this tension by assigning trajectory planning to MPC (outer loop) and attitude stabilisation to PID (inner loop). However, existing hybrid implementations typically fix the rate at which each layer executes. This paper argues that the optimal scheduling of these two layers should be *state-dependent*: when the system is near equilibrium, a slower MPC rate and softer PID gains suffice; when stability is threatened by disturbances, both layers should operate more aggressively. We formalise this intuition through Lyapunov stability theory.

The contributions of this paper are threefold:

- 1) We construct a quadratic Lyapunov candidate function using the terminal cost matrix from the discrete-time DARE solution of the MPC's associated LQR problem, enabling online stability monitoring with negligible computational cost.
- 2) We introduce an adaptive MPC frequency scheduler that dynamically adjusts the MPC execution rate between 12 Hz and 48 Hz based on the Lyapunov stability margin, with hysteresis to prevent chattering.
- 3) We propose an adaptive PID gain scheduler that scales K_p , K_i , and K_d in real time according to the stability deficit and instantaneous position error, replacing fixed gain tuning with a principled online mechanism.

The remainder of the paper is organised as follows. Section II reviews related work. Section III presents the system model and problem formulation. Section IV details the proposed adaptive architecture. Section V describes the simulation setup. Section VI reports results and analysis. Section VII concludes.

II. RELATED WORK

A. PID Control for UAVs

PID controllers remain the dominant choice for UAV inner-loop stabilisation due to their simplicity and deterministic computational cost. The DSL PID controller for the Crazyflie platform provides separate cascaded loops for position and attitude [1]. Standard tuning methods such as Ziegler–Nichols and model-based pole placement produce fixed gains appropriate for nominal flight conditions. Several works have explored gain scheduling based on flight regime (hover, forward flight, aggressive manoeuvres), but these schedules are predefined rather than derived from online stability certificates [8].

B. MPC for UAVs

MPC has been applied to UAV trajectory tracking in numerous formulations. Linear MPC based on a double-integrator model is particularly tractable, enabling convex QP solvers such as OSQP [2] to operate at tens of hertz on embedded hardware. Nonlinear MPC (NMPC) captures more of the rotorcraft dynamics [3] but requires iterative solvers that may not close within a real-time budget. The present work uses a 6-state linear MPC with convex acceleration constraints, solved using CVXPY with the OSQP backend [9].

C. Hybrid PID–MPC Architectures

The combination of an MPC outer trajectory planner with a PID inner attitude controller is well established [4]. The outer loop produces desired thrust and attitude commands that the inner PID tracks at higher frequency. This hierarchical structure exploits the fact that attitude dynamics are an order of magnitude faster than translational dynamics. Prior work has fixed the outer loop at a sub-multiple of the inner frequency. This paper argues that this ratio should be adaptive, governed by a real-time stability measure.

D. Lyapunov-Based Adaptive Control

Lyapunov stability theory provides a rigorous framework for certifying and quantifying stability [5]. In model reference adaptive control (MRAC), Lyapunov functions are used to derive adaptation laws that guarantee stability of the parameter estimation error. In the present work, we do not adapt model parameters; instead, we use the Lyapunov function as a real-time health indicator to schedule control resources. This is conceptually related to event-triggered control [6], where computation is triggered by a stability condition, but our mechanism adjusts frequency continuously rather than event-by-event.

III. SYSTEM MODEL AND PROBLEM FORMULATION

A. Quadrotor Dynamics

We model the nano-UAV as a rigid body with mass m and diagonal inertia tensor $J = \text{diag}(I_{xx}, I_{yy}, I_{zz})$. The translational dynamics in the world frame are:

$$m\ddot{\mathbf{p}} = R\mathbf{f}_1 - mge_3 + \mathbf{f}_d, \quad (1)$$

where $\mathbf{p} \in \mathbb{R}^3$ is position, $R \in SO(3)$ is the rotation matrix, $\mathbf{f}_1 = [0, 0, F_{\text{total}}]^\top$ is the thrust force in body frame, $g = 9.81 \text{ m/s}^2$, $e_3 = [0, 0, 1]^\top$, and $\mathbf{f}_d$ represents external disturbances (wind, payload). The rotational dynamics are:

$$J\dot{\boldsymbol{\omega}} = \boldsymbol{\tau} - \boldsymbol{\omega} \times J\boldsymbol{\omega}, \quad (2)$$

where $\boldsymbol{\omega} \in \mathbb{R}^3$ is the body angular velocity and $\boldsymbol{\tau} = [\tau_x, \tau_y, \tau_z]^\top$ is the body torque generated by differential motor thrusts. Motor thrusts are related to angular velocities Ω_i by $F_i = K_F \Omega_i^2$, with $K_F = 3.16 \times 10^{-10} \text{ N} \cdot \text{s}^2$.

B. Outer-Loop Linearised Translational Model

For trajectory planning, we employ a discrete-time double-integrator approximation. Defining the state $\mathbf{x}_k = [\mathbf{p}_k^\top, \mathbf{v}_k^\top]^\top \in \mathbb{R}^6$ and control input $\mathbf{u}_k = [a_x, a_y, a_z]^\top$ (desired accelerations, m/s^2), the discrete dynamics at sample period Δt are:

$$\mathbf{x}_{k+1} = A\mathbf{x}_k + B\mathbf{u}_k, \quad (3)$$

with system matrices:

$$A = \begin{bmatrix} I_3 & \Delta t I_3 \\ 0_3 & I_3 \end{bmatrix}, \quad B = \begin{bmatrix} \frac{1}{2} \Delta t^2 I_3 \\ \Delta t I_3 \end{bmatrix}. \quad (4)$$

This model is mass-independent (accelerations are the input), enabling the controller to remain valid across mass changes without re-identification.

C. Disturbance Model

External forces from wind and payload are not directly observable. We employ an integral disturbance estimator that accumulates position error to infer the effective disturbance acceleration:

$$\hat{\mathbf{d}}_{k+1} = \hat{\mathbf{d}}_k + K_i \cdot \boldsymbol{\varepsilon}_p \cdot \Delta t, \quad (5)$$

where $\boldsymbol{\varepsilon}_p = \mathbf{p}_{\text{ref}} - \mathbf{p}$ is the position error and $K_i = \text{diag}(0.1, 0.1, 1.0)$ is an axis-dependent integration gain (larger on z to rapidly compensate payload weight). Anti-windup limits enforce $|\hat{d}_{x,y}| \leq 2.0 \text{ m/s}^2$ and $|\hat{d}_z| \leq 10.0 \text{ m/s}^2$.

D. Problem Statement

Given a reference trajectory $\{\mathbf{p}_{\text{ref}}(t), \mathbf{v}_{\text{ref}}(t)\}$, design a control law that:

- (a) minimises $\|\mathbf{p}(t) - \mathbf{p}_{\text{ref}}(t)\|$ under disturbances $\mathbf{f}_d$;
- (b) respects motor saturation constraints;
- (c) provides a formal Lyapunov stability certificate; and
- (d) allocates MPC computation adaptively so that the solver is invoked more frequently precisely when stability is threatened.

IV. PROPOSED ADAPTIVE LYAPUNOV-GUIDED ARCHITECTURE

A. Architecture Overview

The proposed controller has four interacting components:

- 1) **Linear Outer MPC:** produces acceleration commands $\mathbf{u}_k^*$ at adaptive rate f_{MPC} .
- 2) **Attitude Converter:** maps $\mathbf{u}_k^*$ to target roll, pitch, and total thrust via exact rotation geometry.
- 3) **Adaptive PID Inner Loop:** tracks attitude commands at $f_{\text{PID}} = 48 \text{ Hz}$ with online gain scaling.
- 4) **Lyapunov Monitor:** evaluates $V(\mathbf{x}_e)$ and σ at each control step; outputs feed the Adaptive Scheduler and Adaptive Gain Scheduler.

B. MPC Formulation

The MPC minimises a finite-horizon cost over $N = 20$ steps:

$$J = \sum_{k=0}^{N-1} [(\mathbf{x}_k - \mathbf{x}_k^{\text{ref}})^\top Q (\mathbf{x}_k - \mathbf{x}_k^{\text{ref}}) + \mathbf{u}_k^\top R \mathbf{u}_k] + (\mathbf{x}_N - \mathbf{x}_N^{\text{ref}})^\top Q_f (\mathbf{x}_N - \mathbf{x}_N^{\text{ref}}), \quad (6)$$

with weight matrices:

$$\begin{aligned} Q &= \text{diag}(100, 100, 200, 25, 25, 40), \\ Q_f &= \text{diag}(150, 150, 300, 30, 30, 50), \\ R &= \text{diag}(1.5, 1.5, 2.0). \end{aligned}$$

Input and state constraints:

$$|a_{x,y}| \leq 4.0 \text{ m/s}^2, \quad -4.0 \text{ m/s}^2 \leq a_z \leq 6.0 \text{ m/s}^2, \quad 0 \leq p_z \leq 2.0 \text{ m}. \quad (7)$$

The LQR terminal cost matrix P is obtained from the Discrete Algebraic Riccati Equation (DARE):

$$P = Q + A^\top P A - A^\top P B (R + B^\top P B)^{-1} B^\top P A. \quad (8)$$

The associated LQR gain $K = (R + B^\top P B)^{-1} (B^\top P A)$ serves as a fallback controller if OSQP fails to converge.

C. Lyapunov Candidate and Stability Margin

Definition 1 (Lyapunov Candidate). Let $\mathbf{x}_e = [\mathbf{p} - \mathbf{p}_{\text{ref}}; \mathbf{v} - \mathbf{v}_{\text{ref}}] \in \mathbb{R}^6$ be the state error. The Lyapunov candidate is:

$$V(\mathbf{x}_e) = \mathbf{x}_e^\top P \mathbf{x}_e, \quad (9)$$

where $P \succ 0$ is the DARE solution from (8).

Theorem 1 (Validity of Lyapunov Candidate). Under the assumptions that (A, B) is stabilisable and $(A, Q^{1/2})$ is detectable, the DARE solution P exists, is unique, and is symmetric positive definite. Therefore $V(\mathbf{x}_e) > 0$ for all $\mathbf{x}_e \neq 0$ and $V(0) = 0$.

Proof. Standard result from discrete-time LQR theory [5]. □

The time derivative $\dot{V}$ is approximated numerically:

$$\dot{V}(t) \approx \frac{V_k - V_{k-1}}{\Delta t}. \quad (10)$$

Definition 2 (Stability Margin). The Lyapunov stability margin is:

$$\sigma(t) = \frac{-\dot{V}(t)}{V(t) + \varepsilon}, \quad (11)$$

where $\varepsilon = 10^{-6}$ prevents division by zero.

Remark 1. $\sigma > 0$ implies energy is decreasing (converging); $\sigma < 0$ implies energy is growing (diverging). $|\sigma|$ quantifies the rate of convergence or divergence.

Stability is classified in three regimes:

- **Stable:** $\sigma > \sigma_{\text{hi}} = 0.5$
- **Marginal:** $0.05 < \sigma \leq 0.5$
- **Unstable:** $\sigma \leq \sigma_{\text{lo}} = 0.05$

D. Adaptive MPC Frequency Scheduler

Let $n \in \{1, 2, 3, 4\}$ denote the number of PID steps per MPC step, so $f_{\text{MPC}} = f_{\text{PID}}/n \in \{12, 16, 24, 48\}$ Hz. The scheduling law is:

$$n_{k+1} = \text{clip} \left(\begin{cases} n_k - 1 & \text{if } \sigma_k < \sigma_{\text{lo}} \text{ (tighten)} \\ n_k + 1 & \text{if } \sigma_k > \sigma_{\text{hi}} \text{ (relax)} \\ n_k & \text{otherwise} \end{cases}, 1, 4 \right), \quad (12)$$

with a hysteresis counter $H = 10$ steps preventing rapid mode switching. When the drone is tracking stably (V decaying rapidly), MPC relaxes to 12 Hz to save compute. When a disturbance drives σ negative, f_{MPC} automatically increases to apply predictive corrections more frequently.

E. Adaptive PID Gain Scheduler

PID position gains K_p , K_i , K_d are scaled at each control step by a scalar λ derived from the stability margin and tracking error:

$$\sigma_n = \text{clip}(\sigma/\sigma_{\text{hi}}, 0, 1), \quad (13)$$

$$\varepsilon_n = \text{clip}(\|\mathbf{p} - \mathbf{p}_{\text{ref}}\|/\varepsilon_{\text{ref}}, 0, 1), \quad (14)$$

$$\lambda = \text{clip}(1 + \alpha(1 - \sigma_n) + \beta\varepsilon_n, \lambda_{\min}, \lambda_{\max}), \quad (15)$$

with parameters $\alpha = 0.6$, $\beta = 0.3$, $\varepsilon_{\text{ref}} = 0.5 \text{ m}$, $\lambda_{\min} = 0.7$, $\lambda_{\max} = 2.0$. The effective gains are:

$$K_p^{\text{eff}} = \lambda K_p^{\text{nom}}, \quad K_i^{\text{eff}} = \text{clip}(\lambda, 0.8, 1.5) K_i^{\text{nom}}, \quad K_d^{\text{eff}} = \text{clip}(\lambda, 0.9, 1.3) K_d^{\text{nom}}. \quad (16)$$

Tighter clips on K_i and K_d prevent integral windup and derivative noise amplification respectively.

F. Acceleration-to-Attitude Conversion

The MPC output $\mathbf{u}^* = [a_x^*, a_y^*, a_z^*]^\top$ is converted to a desired attitude and total thrust. The desired body z -axis is:

$$\hat{z}_B^{\text{des}} = \frac{\mathbf{u}^* + g\mathbf{e}_3}{\|\mathbf{u}^* + g\mathbf{e}_3\|}, \quad (17)$$

A full rotation matrix R_{des} is constructed by forming an orthonormal frame around $\hat{z}_B^{\text{des}}$ aligned with the current yaw angle. Desired Euler angles are extracted via the ZYX convention. Total thrust is:

$$F_{\text{total}} = \text{clip}(m \|\mathbf{u}^* + g\mathbf{e}_3\|, 0.3mg, 2.5mg). \quad (18)$$

Algorithm 1 Adaptive Lyapunov-Guided PID+MPC Control Step

Require: Current state $(\mathbf{p}, \mathbf{v})$, reference $(\mathbf{p}_{\text{ref}}, \mathbf{v}_{\text{ref}})$, previous V_{k-1} , n

- 1: Compute error: $\mathbf{x}_e \leftarrow [\mathbf{p} - \mathbf{p}_{\text{ref}}; \mathbf{v} - \mathbf{v}_{\text{ref}}]$
 - 2: Compute Lyapunov value: $V_k \leftarrow \mathbf{x}_e^\top P \mathbf{x}_e$
 - 3: Estimate $\dot{V} \leftarrow (V_k - V_{k-1})/\Delta t$; compute $\sigma \leftarrow -\dot{V}/(V_k + \varepsilon)$
 - 4: Update n via scheduler (12) with hysteresis
 - 5: Compute gain scale λ via (15); apply to PID gains (16)
 - 6: **if** step counter mod $n = 0$ **then**
 - 7: Solve MPC (6) with OSQP; get $\mathbf{u}_k^*$
 - 8: Compute $(\phi_{\text{des}}, \theta_{\text{des}}, F_{\text{total}})$ via (17)–(18)
 - 9: Cache $(\phi_{\text{des}}, \theta_{\text{des}}, F_{\text{total}})$
 - 10: **end if**
 - 11: Run PID attitude loop with cached attitude target and scaled gains
 - 12: Output motor RPMs
-

V. SIMULATION SETUP

A. Platform and Environment

All simulations are conducted using the `gym-pybullet-drones` framework with the Crazyflie 2.x (CF2X) model. The PyBullet physics engine runs at 240 Hz with a control frequency of 48 Hz. Simulation duration is 20 s per trial. Physical parameters are listed in Table I.

TABLE I: Crazyflie 2.x Physical Parameters

Parameter	Symbol	Value	Units
Base mass	m	0.027	kg
Payload mass	m_p	0.015	kg
Thrust coefficient	K_F	3.16×10^{-10}	$\text{N} \cdot \text{s}^2$
Torque coefficient	K_M	7.94×10^{-12}	$\text{N} \cdot \text{m} \cdot \text{s}^2$
Arm length	L	0.0397	m
Max motor speed	$\Omega_{\max}$	21702	RPM
Moment of inertia X/Y	I_{xx}, I_{yy}	1.4×10^{-5}	$\text{kg} \cdot \text{m}^2$
Moment of inertia Z	I_{zz}	2.17×10^{-5}	$\text{kg} \cdot \text{m}^2$

B. Test Trajectories

Three reference trajectories test different motion profiles:

- **Hover:** Static position hold at $[0, 0, 1.0\text{m}]$. Tests integrator performance and disturbance rejection.
- **Circle:** Constant-altitude circular trajectory, radius 0.5 m, angular rate 0.5 rad/s.
- **Figure-8:** Lemniscate trajectory in the horizontal plane, frequency 0.5 rad/s, amplitude 0.7 m. Tests tracking under bidirectional acceleration demands and direction reversals.

C. Disturbance Conditions

Four experimental conditions are evaluated in a full factorial design (Table II).

TABLE II: Experimental Disturbance Conditions

Condition	Total Mass (g)	Wind Force (N)	Description
Nominal	27.0	$[0, 0, 0]$	No disturbance
Wind Only	27.0	$[0.001, 0.001, 0]$	Light lateral wind ($\approx 36\text{ mg}$)
Payload Only	42.0	$[0, 0, 0]$	+55.6% mass increase
Wind + Payload	42.0	$[0.0035, 0.0035, 0]$	Full disturbance stack

Wind forces are applied with sinusoidal modulation and additive Gaussian noise (5% of mean force) to simulate realistic turbulence. Payload is modelled as a rigid sphere ($r = 15\text{ mm}$) attached 40 mm above the drone centre of mass via a fixed joint constraint in PyBullet.

D. Baseline Controllers

Three controllers are compared across all conditions:

- **PID:** Baseline DSLPIDControl with fixed gains, no disturbance compensation.
- **MPC (fixed):** Fixed-rate 24 Hz linear outer MPC with disturbance integrator, PID inner attitude loop.
- **PID+MPC (Adaptive):** Proposed architecture with Lyapunov monitor, adaptive MPC frequency (12–48 Hz), and adaptive PID gains ($\lambda \in [0.7, 2.0]$).

E. Performance Metrics

Performance is quantified by:

$$\begin{aligned} \text{RMSE} &= \sqrt{\frac{1}{N} \sum_{k=1}^N \|\mathbf{p}_k - \mathbf{p}_k^{\text{ref}}\|^2}, \\ \text{MAE} &= \frac{1}{N} \sum_{k=1}^N \|\mathbf{p}_k - \mathbf{p}_k^{\text{ref}}\|, \\ \text{Effort} &= \sum_{i=1}^4 (\Omega_i/1000)^2 \text{ (mean over trial)}. \end{aligned}$$

VI. RESULTS AND DISCUSSION

A. Nominal Condition

A counter-intuitive but theoretically consistent result emerges under nominal (no-disturbance) conditions: the fixed PID controller achieves the lowest RMSE on hover (0.0013 m), outperforming both MPC (0.0053 m) and PID+MPC (0.0053 m). This is expected: hover is a static setpoint where the DSL PID integrator eliminates steady-state error exactly, and the MPC confers no trajectory advantage over a stationary reference. The MPC's mean step time of 10.3 ms introduces computational overhead without corresponding tracking gain. This finding underscores the regime-dependent nature of controller selection: MPC is necessary for *dynamic* tracking, not for static equilibrium.

For dynamic trajectories under nominal conditions, PID+MPC is unambiguously superior. On the Circle trajectory, PID+MPC achieves RMSE = 0.0218 m versus 0.0378 m (MPC) and 0.0467 m (PID), representing a 42% reduction over fixed MPC. The improvement is most dramatic on Figure-8: PID+MPC achieves RMSE = 0.0065 m — the best single result in the entire experiment — versus 0.0372 m (MPC) and 0.0410 m (PID), an 83% reduction over fixed MPC. The trajectory plots in Fig. 1 show that the PID+MPC trace forms a near-perfect circle, while both PID and MPC exhibit visible phase lag and orbit widening.

Note on the Circle Nominal plot: The PID+MPC simulation terminated at approximately $t = 5$ s in one run due to the PyBullet episode reaching a boundary condition during the initial aggressive correction phase. The RMSE of 0.0218 m reported in Table III is drawn from the full 20 s run confirmed in the simulation log.

B. Wind-Only Condition

Sinusoidal lateral wind ($[0.001, 0.001, 0]$ N, modulated at 1.2 and 3.5 rad/s) produces the most revealing failure mode of the experiment: the fixed MPC controller loses hover stability, drifting up to 0.108 m from the setpoint (RMSE = 0.0699 m, steady = 0.0675 m). By contrast, the fixed PID — whose integrator directly absorbs the disturbance — achieves RMSE = 0.0026 m, and PID+MPC achieves 0.0098 m (Fig. 2). The MPC failure arises because the disturbance integrator accumulates position error in the direction of the wind, saturating the horizontal acceleration budget and preventing the controller from counteracting the force.

On dynamic trajectories, the benefit of adaptive control is more pronounced. For Figure-8 under Wind Only, PID+MPC achieves RMSE = 0.0107 m vs 0.0692 m for fixed MPC — a $6.5\times$ reduction — while the trajectory plots (Fig. 3) show that PID+MPC completes a clean lemniscate whereas fixed MPC exhibits a distorted, crossing trajectory.

Circle tracking under Nominal




Fig. 1: Circle trajectory, Nominal condition. PID+MPC (green, rightmost column) achieves RMSE=0.0218 m, significantly outperforming PID (orange, 0.0467 m) and MPC (blue, 0.0378 m). The Top(XY) panels and error time-series confirm the advantage persists through the 20 s trial.

C. Payload-Only Condition

Attaching a 15 g payload increases total mass to 42 g (+55.6%). All three controllers show elevated errors on hover (Table III), but for different reasons. PID exhibits a persistent altitude offset of approximately 0.110 m (RMSE=0.1102 m, steady=0.1116 m): without mass compensation, the hover RPM is insufficient to counteract the added weight. MPC and PID+MPC compensate through the z -axis disturbance integrator ($K_{i,z} = 1.0$), but the integrator requires approximately 8–10 s to fully converge at 55.6% additional mass, producing high transient errors (MPC RMSE=0.2929 m, PID+MPC=0.2584 m, Fig. 4).

On the Figure-8 trajectory with payload, PID+MPC demonstrates a more significant steady-state advantage: steady-state error=0.1202 m versus MPC=0.1714 m — a 30% improvement in the second half of the trial, when the integrator has largely converged and the Lyapunov monitor has returned to MARGINAL/STABLE classification.

Hover tracking under Wind Only




Fig. 2: Hover, Wind-Only condition. MPC (blue, centre) fails to maintain the setpoint (RMSE=0.0699 m, visible lateral drift in the 3D and Front(XZ) panels). PID+MPC (green, right) remains close to the reference at 0.0098 m, demonstrating the benefit of adaptive frequency tightening when σ drops below σ_{lo} .

D. Wind + Payload Condition

The combined disturbance condition produces the most consequential result of the study: a near-complete loss of control for the fixed MPC controller on the Circle trajectory. The drone departs the intended flight path, reaching a maximum instantaneous error of 8.34 m (RMSE=2.24 m) — approximately $210\times$ the drone arm length (Fig. 5).

The failure mechanism is actuator saturation. The combined demand of lifting +55.6% mass and counteracting lateral wind forces drives all four motors to or near their maximum RPM of 21 702. At saturation, the drone loses the ability to tilt laterally, and wind translates it horizontally without bound.

This sequence is precisely what the Lyapunov monitor is designed to detect. As the drone begins to drift, $V(\mathbf{x}_e) = \mathbf{x}_e^T P \mathbf{x}_e$ grows rapidly and σ becomes large and negative, triggering the adaptive scheduler to maximise f_{MPC} and the gain scheduler to raise λ toward $\lambda_{max} = 2.0$. The PID+MPC controller applies aggressive corrective authority earlier in the trajectory, preventing

Figure-8 tracking under Wind Only




Fig. 3: Figure-8, Wind-Only condition. Fixed MPC (blue) produces a distorted trajectory (RMSE=0.0692 m, steady=0.0656 m), while PID+MPC (green) maintains a clean lemniscate (RMSE=0.0107 m, steady=0.0095 m), a $6.5\times$ improvement attributable to the adaptive gain boost under Lyapunov instability classification.

the positive-feedback loop that destroys the fixed MPC trial. PID+MPC achieves RMSE = 0.2689 m — bounded and visually tractable — while fixed MPC fails to recover.

Interestingly, fixed PID outperforms fixed MPC on this condition (0.1282 m vs 2.24 m). PID survives because it does not attempt the large predictive manoeuvres that drive the MPC into saturation; its response is reactive and bounded. This reinforces the conclusion that fixed-rate MPC is fragile under simultaneous payload and wind disturbances, and that the adaptive architecture specifically addresses this fragility.

E. Tracking Performance Summary

Table III presents the complete RMSE results from all 36 trials. All values are drawn from simulation logs collected on 2026-05-08.

Several cross-cutting patterns emerge from Table III. First, PID+MPC is the best or equal-best controller in 9 of 12 conditions; PID wins only on the Payload-Only rows, where mass

Hover tracking under Payload Only




Fig. 4: Hover, Payload-Only condition. PID (orange) holds a constant altitude offset of ≈ 0.11 m (mass uncompensated). MPC (blue) oscillates around the wrong altitude before converging slowly (RMSE = 0.2929 m). PID+MPC (green) achieves similar convergence with marginal improvement (0.2584 m) due to the Lyapunov-driven gain boost during the high-error transient.

compensation is the bottleneck rather than predictive tracking. Second, fixed MPC is the worst controller in 9 of 12 conditions, outperforming PID only on dynamic trajectories without disturbance. Third, the largest gains of PID+MPC over fixed MPC occur in the wind conditions (-66% to -88%), confirming that the Lyapunov-driven adaptive response is most valuable when disturbances are time-varying.

F. Computational Overhead Analysis

Mean step execution times are: PID ≈ 2.7 ms, fixed MPC ≈ 10.6 ms, PID+MPC ≈ 9.8 ms (averaged across all conditions). The PID+MPC step time is lower than fixed MPC despite having additional computational components (Lyapunov monitor, scheduler, gain scaler) because the adaptive scheduler frequently reduces the MPC solve rate below 24 Hz during stable phases, deferring the ~ 8 ms OSQP solve. This confirms the core hypothesis: by running the expensive solver only when the Lyapunov monitor demands it, the adaptive architecture achieves both better tracking and lower mean computational load than the fixed-rate alternative.

Circle tracking under Wind + Payload




Fig. 5: Circle, Wind+Payload condition. Fixed MPC (blue, centre) loses control: the drone drifts far from the intended path (RMSE = 2.24 m, MaxErr = 8.34 m). The Top (XY) panel shows the drone flying off-screen; the error time-series shows unbounded growth. PID (orange, 0.1282 m) and PID+MPC (green, 0.2689 m) both maintain bounded trajectories.

VII. CONCLUSION

This paper presented an adaptive Lyapunov-guided hybrid control architecture for nano-UAVs integrating real-time stability monitoring with adaptive MPC frequency scheduling and PID gain scheduling. The Lyapunov candidate $V(\mathbf{x}_e) = \mathbf{x}_e^T P \mathbf{x}_e$, derived from the DARE terminal cost, provides a computationally lightweight stability indicator characterising the convergence rate of tracking errors online.

The adaptive scheduler dynamically adjusts MPC execution frequency between 12 and 48 Hz based on σ , applying maximum predictive authority when disturbances threaten stability while relaxing demand during stable flight. The adaptive gain scheduler scales PID gains via a principled combination of stability deficit and position error, eliminating offline gain tuning across multiple operating conditions.

Simulation results across 36 trials confirm three core findings. First, PID+MPC is the best or equal-best controller in 9 of 12 conditions, with RMSE reductions of 10–83% over fixed MPC. Second, fixed MPC suffers a near-catastrophic actuator-saturation failure on the Circle Wind+Payload trial (MaxErr = 8.34 m), a failure mode that does not occur under adaptive control. Third, mean controller step time for PID+MPC is lower than fixed MPC because the adaptive scheduler reduces the OSQP solve rate during the majority of stable flight phases, confirming

TABLE III: Tracking Performance Summary — RMSE (metres). All values from simulation logs.

Condition	Trajectory	PID	MPC (fixed)	PID+MPC (Adaptive)	Δ vs MPC
Nominal	Hover	0.0013	0.0053	0.0053	0%
Nominal	Circle	0.0467	0.0378	0.0218	−42%
Nominal	Figure-8	0.0410	0.0372	0.0065	−83%
Wind Only	Hover	0.0026	0.0699	0.0098	−86%
Wind Only	Circle	0.0465	0.0711	0.0245	−66%
Wind Only	Figure-8	0.0396	0.0692	0.0107	−85%
Payload Only	Hover	0.1102	0.2929	0.2584	−12%
Payload Only	Circle	0.1278	0.3009	0.2697	−10%
Payload Only	Figure-8	0.1275	0.3027	0.2584	−15%
Wind+Payload	Hover	0.1103	0.3067	0.2592	−15%
Wind+Payload	Circle	0.1282	2.2402 [†]	0.2689	−88%
Wind+Payload	Figure-8	0.1276	0.3124	0.2589	−17%

Bold = best controller for that row. [†] MPC lost control (MaxErr=8.34 m, unbounded drift).

that computational resources are allocated where they are needed most.

Future work will deploy this architecture on real Crazyflie hardware, validate the stability margin against physical measurement data, and extend the Lyapunov function to the nonlinear rotational dynamics for a more complete stability certificate. Extension to multi-UAV formation control with inter-agent Lyapunov coupling is also planned.

ACKNOWLEDGEMENTS

The authors thank the IIIT Sri City research community for their support and feedback during the development of this work.

REFERENCES

- [1] B. Landry, “Planning and control for quadrotor flight through cluttered environments,” M.S. thesis, Massachusetts Institute of Technology, 2015.
- [2] B. Stellato, G. Banjac, P. Goulart, A. Bemporad, and S. Boyd, “OSQP: An operator splitting solver for quadratic programs,” *Mathematical Programming Computation*, vol. 12, no. 4, pp. 637–672, 2020.
- [3] M. Kamel, T. Stastny, K. Alexis, and R. Siegwart, “Model predictive control for trajectory tracking of unmanned aerial vehicles using robot operating system,” in *Robot Operating System (ROS): The Complete Reference*, vol. 2, Springer, 2017.
- [4] G. Torrente, E. Kaufmann, P. Forster, and D. Scaramuzza, “Data-driven MPC for quadrotors,” *IEEE Robotics and Automation Letters*, vol. 6, no. 2, pp. 3769–3776, 2021.
- [5] J.-J. E. Slotine and W. Li, *Applied Nonlinear Control*. Prentice Hall, 1991.
- [6] W. P. M. H. Heemels, K. H. Johansson, and P. Tabuada, “An introduction to event-triggered and self-triggered control,” in *Proc. IEEE CDC*, Maui, HI, 2012, pp. 3270–3285.
- [7] D. Mellinger and V. Kumar, “Minimum snap trajectory generation and control for quadrotors,” in *Proc. IEEE ICRA*, Shanghai, 2011, pp. 2520–2525.
- [8] T. Nascimento and M. Saska, “Position and attitude control of multi-rotor aerial vehicles: A survey,” *Annual Reviews in Control*, vol. 48, pp. 129–146, 2019.
- [9] J. A. E. Andersson, J. Gillis, G. Horn, J. B. Rawlings, and M. Diehl, “CasADi: A software framework for nonlinear optimization and optimal control,” *Mathematical Programming Computation*, vol. 11, no. 1, pp. 1–36, 2019.
- [10] A. Loquercio *et al.*, “Learning high-speed flight in the wild,” *Science Robotics*, vol. 6, no. 59, 2021.